

ClaudeAssist: A Secure MCP-Based AI Workspace Manager

Sace Singh
 CSE Department
 Amity University Uttar Pradesh
 Noida, India
 airforceschool2008@gmail.com

Armaan Marya
 CSE Department
 Amity University Uttar Pradesh
 Noida, India
 armaanmarya2004@gmail.com

Sanjeev Tomar
 CSE Department
 Amity University Uttar Pradesh
 Noida, India
 skumar8@amity.edu

Vikas Kamra
 School of Computer Science and
 Engineering, IILM University
 Greater Noida, India
 kamra1984@gmail.com

Abstract—With the ongoing growth of digital academic and professional work processes, massively huge amounts of documents, emails, and research data are being created on local and cloud-provided platforms, thus causing disorganized information management and decreased productivity. This paper proposes a full system of an intelligent (AND privacy-sensitive) workspace assistant synthesizing the usage of large language models reasoning with individual computing positions. The proposed system, ClaudeAssist, uses a local-first system that is built on the Model Context Protocol (MCP) to support secure interaction between an LLM and local and cloud resources. State-of-the-art contextual retrieval and orchestration algorithms are engulfed with the help of LangChain and LlamaIndex to improve the level of reasoning and persistent memory of heterogeneous sources of data. Observations obtained during experiments prove that the local contextual reasoning together with the automated execution of the tools is more reliable and privacy-aware automation than the traditional cloud-based AI assistants. It is a system that is able to effectively organize documents, summarize them and even automate emailing without losing the ultimate user control obtained through sandboxed execution and versioned checkpoints. The results suggest that workspace management based on LLM can be safely scaled without sensitive information disclosed and allows creating a platform of trustful, autonomous, and privacy-oriented AI-based help in the sphere of academia and various professional spheres.

Keywords—Claude LLM, Model Context Protocol, AI Workspace Manager, Google Drive Integration, LangChain.

I. INTRODUCTION

The fast growth of digital academic and professional ecosystems has presented enormous challenges in the handling of huge amounts of documents, communications, and research artifacts that are shared with the different platforms. Complex digital settings are created by local storage systems, cloud repositories, and email services that usually divide valuable contextual information and use it inadequately. Contemporary patterns of interaction are more dependent on unstructured textual data implicitly reflecting the workflow intent, task priorities, and knowledge development outside of users [1].

In contrast to conventional methods of productivity that rely on a manual system, and discrete software programs, intelligent automation provides ongoing, contextual

workspace support. Nevertheless, current solutions have significant limitations concerning privacy of data, disjointed integration, scalability and reliance on centralized cloud systems or data layer hence the inability to be used in sensitive academic or research processes [2]. ClaudeAssist is an intelligent system architecture, refer Figure 1.




Fig. 1. ClaudeAssist System Architecture

The recent developments in the field of natural language processing have greatly enhanced the skills in analyzing unstructured textual information created in digital working environments and provided a more profound insight into the intent of a user, its contextual significance and knowledge tendencies [3]. Upgrade of these systems is evidenced in this research work. Intelligent systems that operate under a text-based interface are often based on classification, retrieval, reasoning frameworks to make inferences about relevance to tasks and automate decisions made out of fragmented interactions, i.e., documents, emails, and logs [4].

Modern deep learning architectures that have the potential to learn rich semantic representations have had an evolutionary path starting with primitive rule-based and classical machine learning methods that used handcrafted features to prevalent systems helpful in semantic research such as unsupervised clustering, semi-supervised learning, generative adversarial models, and reinforcement learning. The first solutions had used methods like support vector machines and probabilistic models of document categorization and the current solutions are starting to use transformer-based architectures to understand the context [5].

These new models are more equipped to follow long range dependencies, implicit intention and semantic associations

that exist between heterogeneous data sources and are necessary in achieving the accurate and reliable automation of intelligent workspace [6]. Old-fashioned intelligent workspace fundamentals like disjointed sources of data, uneven labelling of context, ethical issues of automation with respect to privacy, and constraints of translucent model understanding can be represented. The changes in ways of linguistic expression, diversity of user behaviours, and constraints presented by a certain platform, the heterogeneity of document formats also complicate the creation of universally adaptable systems, refer Figure 2.




Fig. 2. Evolution of Intelligent Workspace Technologies

In this paper, an intelligent workspace architecture is provided in detail which manages and reasons on distributed local and cloud-based data using a privacy-first architecture. This approach should be proposed to focus on formal contextual modelling, reliable tool-based automation and attentive system design in order to maintain transparency and ethical usage [7]. Combining state of the art language model reasoning with artificial intelligence, controlled execution and persistent memory, the system is set to provide trustworthy, convenient, and effective digital support at the expense of preserving the confidence of users and eliminating risks that may arise with the sensitive data processing [8].

II. LITERATURE REVIEW

A. Privacy-Preserving Intelligent Workspace Systems

Initial privacy-conscious workspace systems involved cloud automation and assistants based on rules to control documents, emails, and workflows. Satyanarayana et al. emphasized edge computing in matters of diminishing the reliance on clouds, whereas Greenwood et al. emphasized the local processing of personal data securely [9]. Subsequently, structured interactions were added by the tool-augmented language model frameworks, which restricted file, memory, and system tool access by the model. The majority of strategies are still cloud-based, which compromise the ability of users to control sensitive data [10].

Lightweight standards such as the Model Context Protocol (MCP) of Anthropic allow sandboxed, schema-driven communication, allowing unified local-cloud integration and stable contextual memory, driving privacy-first workspace managers such as ClaudeAssist.

B. Large Language Models for Contextual Automation

1. New developments within the large language models allow ClaudeAssist to computerize the contextual workflow that incorporates resources both locally and in the cloud, and retains sensitive information. Key architecture is:
2. Transformer-based LLMs have the ability to learn contextual and sequential patterns in raw text or structured workflow; can be further extended to external sources of knowledge, system tools, or local-cloud integration. Memory modules are effective in keeping a

history of tasks and also allow continuity of understanding current workflows between sessions.

3. ClaudeAssist has attention and hierarchical attention mechanisms, which enable the assist to be focused on salient documents, emails or workflow segments. The study by researchers [11] displayed the same structured anthropomorphic network in order to make the tool-augmented LLMs prioritize the necessary data and filter the availability of external resources.
4. Memory-augmented architectures create task and user contact dependencies that facilitate understanding of context in the long run, flawless local-cloud functionality, and the capacity to retrieve previous choices or user interests without sensitive information disclosure.
5. Flexible design enables the addition of external tools, APIs or data bases, enabling orderly automation and privacy protection, sandbox provided execution, and frequent access to both local and distant resources to reduce the risk of data leakage [12].

C. Hybrid Local-Cloud AI Architectures

It is important to distinguish the tasks of local resources and cloud resources in order to implement the hybrid local cloud AI system and guarantee its efficiency. Common approaches include:

1. Local device calculations which process sensitive user information and working context temporary compute information, provides privacy with minimal latency to real-time tasks processing.
2. Hybrid task orchestration: Since certain AI computations are assigned to the local devices or cloud servers, a dilemma exists between the privacy of the user data and the efficiency in the computations. Local execution means that sensitive contextual information and workflow details are not lost, and the cloud execution gives the opportunity to update many models, have a persistent contextual memory, and includes large datasets to enhance the automation [13]. A way to balance these aspects is by providing strategies to optimize tasks placement, impose sandboxed processing and retain a flowed interaction of local and remote environments which privacy conscious, responsive, and contextually sensitive workspace management.
3. No longer having to maintain a contextual memory across sessions due to the combination of cloud-based processing, which is used in heavy calculations, updating model and large-scale data integration and is scalable [14].




Fig. 3. Hybrid Local-Cloud AI Architecture Diagram

The best data set problems are the annotation bias, the imbalance of classes and the instances where the public benchmarks are limited in quantities [15]. Furthermore, the majority of data sets are focused on the English language that means that is necessary to preprocess it more language-specifically and to re-train the model to consider and underrepresented languages, refer Figure 3.

III. RESEARCH METHODOLOGY

A. MCP Server and Local Workspace Integration

In order to create ClaudeAssist as a safe and smart local workspace manager, one takes the following steps:

1) Local Environment Connection

The local files of the user are linked to Claude through local MCP server. It supports multi-format (.txt, .docx, .pdf, .md), and the sandboxed read/write access in order to guarantee privacy. Logging is done on all the operations and having versioned checkpoint allows one to roll back, audit and be able to reproduce edits, summaries and automated actions.

2) Tool Layer Automation

Custom MCP tools do allow independent operations like summarization, keyword extraction, and search of files as well as code generation. All tools interact with each other using organized MCP schemas, with external calls limited and processing with privacy as its first priority [16]. Multi-agent orchestration enables Claude to achieve coordination with tool agents when doing sophisticated local workflows.

3) Contextual Memory and Logging

Local edits, summaries, and task metadata data are embedded in lightweight databases (JSON/SQLite). The embeddings of vectors offer the ability to have consistent contextual memory, facilitating retrieval-enhanced reasoning and keeping the data on-a-chip, refer Figure 4.




Fig. 4. Enhanced Context and Reasoning Diagram

B. Cloud Integration and Autonomous Workspace Management

Taking Claude Assist as the system to the physical scope of local files, the system incorporates cloud resources and authenticates independent multi-source reasoning:

1) Connection

Google Drive and email accounts (Gmail/ Outlook) are secured through the OAuth 2.0. Metadata of the files resided locally and provides context-sensitive search, summarization as well as offline support [17]. There is a summative view of

emails, ranking of emails, and the creation of draft responses with the help of the stored context without breaking local first privacies.

2) Tools

The secure access to Google Drive as well as email accounts (Gmail/Outlook) is defined by OAuth 2.0. Metadata is stored locally to provide context-aware search of files, summarization and offline access. Priorities and curl summaries of emails are generated; draft responses are produced using stored context and local-first privacy is preserved [18].

3) Memory

Promptly workflow mapping and dependency visualization (with Cytoscape.js) is offered as optional modules. Every step is monitored, recorded and determined as a checkpoint to guarantee a transparency, reproducibility, and privacy depending on local and cloud-integrated tasks [19].

IV. FEATURE ENGINEERING AND CONTEXT MODELING

In order to transform raw user interactions and documents into meaningful model inputs, one has to extract features of structural, semantic, operational as well as contextual patterns of workspace activity [20]. Below are the key strategies:

1. Structural Features

- Document Patterns: Sentences, bullet points, headings, and code blocks are used to show how the text or the file has been organized and the purpose of that file.
- File Metrics: Word count, sentence length, and content density are indicators of complexity and length of the tasks or the project.

2. Semantic and Topic Features

- Embeddings: Word2Vec, GloVe or contextualized embeddings (BERT/Claude embeddings) captures the semantic relations between the instructions, emails or research content; document embeddings summarize files, and turns it into a single vector [21].
- Topic Modeling - LDA or NMF can find latent themes across the documents, emails, or project notes; prevalence of topics reflects priorities in workflow and relative relevance.

3. Operational / Interaction Features

- Types of actions: The method of edits, summarization, or file access will indicate the involvement of the user and progress in the relevant task.
- Timing & sequences: The time that has been spent on a file, the periods between the actions, or the sequence of the tasks show the dynamics in the workflow [22].
- Tool Usage: Number and style of invocation MCP tool (e.g., summarizer, search, email generator) reflect the reliance on automation and also task complexity.

4. Statistical Reasoning and Structural Pathways

- N-grams: Routine tasks or talked about ideas are represented into document, email, and code files as common phrases.

- TF-IDF: To weight words by specificity to documents or projects in order to highlight unique instructions or context-sensitive instructions.
- Length Metrics: File sizes, number of edits or words/sentences is the indicator of workload and document complexity.

5. Domain & Auxiliary Features

- File Metadata: Timestamps, authorship, version history, and file type provide information related to lifecycle and relevancy of the content in a workspace [23].
- Collaboration Metrics: Frequency of email exchanges, shared documents and team interactions is a measure of engagement and dependencies in workflow.

Structural, semantic, behavioral features help to combine the context and reasoning of Claude Assist.

V. PROPOSED SYSTEM ARCHITECTURE

It requires a sensible system architecture in order to support safe, context-aware and privacy-enforcing automation using local-first intelligent workspace execution.

A. A Model Context Protocol (MCP) - Based Architecture

- Claude large language model is combined with the help of the Model Context Protocol that specifies the structured schemes of the tool usage, the commercialisation of the context between the local and the cloud resources and the admission based on permission. MCP allows the model to execute restricted write, read, and reason operation on documents, emails, and metadata without unveiling the system, hence preservation of data privacy.
- There is a layered orchestration approach that is used to control intricate workflows within local files and cloud services. In the former, documents or emails are handled at the first level in sandbox deductive tools. Contextual aggregation (at the second level) is an aggregation of related resources in order to facilitate task-level reasoning like summarization or prioritization. In the third tier, memories compose enduring memories that incorporate historical interactions and embeddings and feed on these to produce unified workspace learning and continuity between sessions whilst preserving versioned checkpoints and rollback.

B. Contextual Reasoning and Memory-Oriented Architectures

Context-aware architectures facilitate continuity between interactions across interactions, and allow intelligent workspace systems to model sequential information. Persistent contextual memory and attention-based reasoning is used in ClaudeAssist to track the changing documents, emails and tasks over time.

- Sequential workspace activities are modelled by a contextual memory mechanism that will trace the evolution of documents, tasks and interactions over lengthy time periods.
- Embedding-based representations represent the gradual extent of any user files and communications

and thus this allows persistent contextual knowledge across sessions and various workspace assets.

- Attention driven reasoning emphasizes the most pertinent documents, messages or metadata needed to execute tasks efficiently and accurately and in context.
- Temporal context aggregation enables the system to change responses and automation strategies depending on past interactions, revisions and changing user workflow.

C. Central Motive Engine of Reasoning Transformers

The main intelligence layer of ClaudeAssist is transformer-based reasoning, which allows deep contextual knowledge of the workspace data of various types. The architecture uses large language models to run documents, emails, and metadata concurrently and enables scalable and consistent automation of local and cloud environments.

- Workspace tokens can be processed in parallel with the self-attention mechanism, which helps the system to capture the immediate file-level context and long-range dependencies between projects and communication histories.
- Pre-trained transformer models are task-oriented, making it possible to do accurate summary, classification and decision support without necessary retraining of models on a task-specific basis.
- The model has contextual adaptation mechanisms, which enable it to adapt its reasoning to academic, research and professional domains, using user specific data and history of interaction in a privacy-sensitive way.

D. Tool-Oriented Modular Execution Framework

The systems architecture proposed has employed a system of modular execution, which allows organizing and safe automation of tasks using specialized tools. Each tool has the purpose of executing a particular operation like file handling, document overview, email or retrieving a context, and the tool works within a limited set of permissions spelt out by the Model Context Protocol.

- Tool modules run in sandboxed environment and grant limited access to system resources and prevent unwanted data leakage on automated activities.
- Task orchestration systems bring together a variety of tools to facilitate the complex workflows to enable sequential and parallel execution to respond to the contexts based on the situations.
- Modular extensibility makes the integration of new tools to be able to easily be implemented without having to jeopardize the security or performance of the system.

All in all, the suggested architecture is a combination of secure local implementation, a logic in the context, and modular automation that can provide scalable and privacy-sensitive intelligent workspace system.

VI. IMPLEMENTATION AND INTERFACE

ClaudeAssist was deployed and certified with a structured interaction-based evaluation framework on the scheme of smart workspace utilization records.

A. Preparation of System Interaction Data

- **Source / Logging:** Real-time ClaudeAssist session traces (tool calls, executions of MCPs, and workspace commands) were automatically logged in a local file to restrict privacy and contextual integrity.
- **Preprocessing:** Command prompts, system responses and execution metadata were normalized and tokenized and place value-aligned. The organized states of interaction were converted into model-usable symbolizations to be analyzed and coordinated further.

B. Implementation Setup

- **Component Choice:** The system is built around MCP based components including filesystem access, calculators and simulation of the virtual machine, which is controlled by a central conversational controller.
- **Implementation and evaluation:** Interaction: Interactions were experimented under repeated task conditions to test interfaces with regard to reliable tool invocation, contextual persisting and reliable response generation under different workflows.
- **Robustness Handling:** Permission scopes were validated; unauthorized access was controlled and failures in the implementation of the execution were handled gracefully within the interface.




Fig. 5. ClaudeAssist Interface

C. Evaluation Metrics

- **Usability:** Work precision and accuracy of performed actions.
- **Responsiveness:** System responses to multi-step workflow are timely and coherent.
- **Context Retention:** Capacity to save and correct between-session conversation and operational context.

- **Interface Reliability:** connector stability, MCP tools stability, and local execution stability without data leakage or data loss of context.

The current paper assesses the options of ClaudeAssist, contrasting its local, cloud-integrated and multi-agent architectures regarding automation of tasks and contextual reasoning, refer Figure 5.

A comparison of the performance comparison between ClaudeAssist and two control settings: Local-only and Cloud-Integrated is listed in Table I. According to the Local-only setup, there is an accuracy of 73.2, precision of 72.8, recall of 72.5, F1-score of 72.6 and the AUC-ROC of 0.78 which shows the ability to perform simple tasks and little contextual processing.

Using cloud storage, email and document access, Cloud-Integrated set up delivers higher performance of 81.0% accuracy, 80.7% precision, 80.5% recall, 80.6% F1-score, and 0.85 AUC-ROC. The full-fledged multi-agent configuration of ClaudeAssist has the best metrics of 89.7% accuracy, 89.4% precision, 89.1% recall, 89.2% F1-score and 0.93 AUC-ROC and best multi-source reasoning, automated workflow management, persistent contextual memory and a perfect combination of local and cloud environments contribute to its being the most effective and reliable configuration in the end.

TABLE I: PERFORMANCE COMPARISON OF PROPOSED METHOD WITH BASE METHODS

Model Configuration	Accuracy %	Precision %	Recall %	F1-Score %	AUC-ROC
Local-Only	73.2	72.8	72.5	72.6	0.78
Cloud-Integrated	81.0	80.7	80.5	80.6	0.85
ClaudeAssist (Multi-Agent)	89.7	89.4	89.2	89.2	0.93

It is clear that the performance increases with every layer of integration, as reflected in this table with the best contextual reasoning, workflow automation, and privacy-aware task management being obtained with the full multi-agent ClaudeAssist setup.

VII. EXPERIMENTAL ANALYSIS AND DISCUSSION

A. Interpretability

Deep neural models such as RNNs, CNNs, and transformers are opaque models but precise. There must be transparency to develop trust and facilitate informed choices. Key strategies include:

- **Significance of features:** Comparison of the contribution of words or features to predictions.
- **Attention visualization:** Circle tokens or sentences model is attending them.
- **Model-agnostic explanations (e.g. LIME, SHAP):** Produce a human understandable explanation of output.
- **Rule-based summaries:** Incorporate neural outputs with theory-based, explicit rules.

The methods are useful in examining mistakes, a sense of responsibility, and in leading to the effective decision-making.

B. Ethical Considerations

Ethical compliance must exist in AI workspace and automation:

- Privacy: De-identify and store sensitive data in a safe manner.
- Consent: Get overt user permission concerning use of data.
- Fairness: Spot biases and have equal, diverse representation.
- Annotation The quality: Make sure that labels are standardized and correct.
- Human Control: AI does not negate human judgments, but protocols in case of essential situations should be present.
- Transparency & Governance: Report capabilities, risks, and comply with ethics boards and law.

Interpretability combined with ethical safeguards ensures ClaudeAssist operates reliably, responsibly, and trustworthily.

VIII. CONCLUSION AND FUTURE WORK

ClaudeAssist is a new generation personal work assistant, which combines Claude LLM with a local-first, privacy-protecting MCP architecture and cloud infrastructure. ClaudeAssist supports autonomous summarization, document organization, email management, code assistance, and research assistance by relating local files, emails and cloud storage to multi-agent reasoning and workflow with tool augmentation. The system has a persistent versioned checkpoint based on a vector based memory, attention driven reasoning and the capability of doing contextual understanding correctly and reproducible actions which makes it efficient and reliable. The comparison of ClaudeAssist with local-only and cloud-integrated architectures indicates that the full multi-agent architecture of ClaudeAssist is more accurate, precise, recalls higher, more F1-score, and AUC-ROC, validating the advantages of having hybrid structures and advanced automation systems.

The next round of effort will be to empower the assistant in a range of more productivity platforms including calendar, task management, and messaging API applications that can allow smooth workflow automation of other platforms. Voice interaction, and natural language control will be available in an offline mode to offer a more intuitive way of performing tasks. The advanced visualization modules will integrate user-monitoring the workflow dependencies, contextual memory and tool interaction in real-time. Lastly, additional improvements in ethical protection, privacy, and human-in-the-loop control will be the aspects guaranteeing that ClaudeAssist will be secure, trusted, and adjustable to various academic, research, and professional settings to become a full-fledged assistant in managing intelligent digital workspace.

ACKNOWLEDGMENT

The authors heartly thank the organizers for their support.

REFERENCES

- [1] W. Zhou, X. Zhu, Q.-L. Han, L. Li, X. Chen, S. Wen, and Y. Xiang, "The security of using large language models: A survey with emphasis on ChatGPT," *IEEE/CAA Journal of Automatica Sinica*, vol. 12, no. 1, pp. 1–26, Jan. 2025, doi: 10.1109/JAS.2024.124983.
- [2] M. Bezzi, "Large language models and security," *IEEE Security & Privacy*, vol. 22, no. 2, pp. 2–10, Mar.–Apr. 2024, doi: 10.1109/MSEC.2023.3345568.
- [3] Y. Liu, P. Li, J. Chen, and X. Zhang, "Trustworthy large language models: Foundations and challenges," *IEEE Access*, vol. 12, pp. 145321–145339, 2024, doi: 10.1109/ACCESS.2024.3378912.
- [4] L. Shokri, M. Stronati, C. Song, and V. Shmatikov, "Membership inference attacks against machine learning models," *IEEE Transactions on Information Forensics and Security*, vol. 15, pp. 255–270, 2020, doi: 10.1109/TIFS.2019.2934689.
- [5] R. Mireshghallah, A. Taram, P. Papernot, and M. Jaggi, "Privacy in large language models: Attacks, defenses, and future directions," *IEEE Security & Privacy*, vol. 21, no. 6, pp. 56–66, Nov.–Dec. 2023.
- [6] S. Wen, W. Zhou, J. Zhang, Y. Xiang, and W. Zhou, "Adversarial attacks and defenses on deep learning models," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 32, no. 2, pp. 521–538, Feb. 2021.
- [7] H. Xu, W. Han, and X. Li, "Explainable artificial intelligence in deep learning: A survey," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 34, no. 5, pp. 1965–1982, May 2023.
- [8] V. Kamra, P. Kumar and M. Mohammadian, "Formulation of an Elegant Diagnostic Approach for an Intelligent Disease Recommendation System," 2019 9th International Conference on Cloud Computing, Data Science & Engineering (Confluence), Noida, India, 2019, pp. 278–281.
- [9] Q. Liu, P. Li, W. Zhao, W. Cai, S. Yu, and V. C. M. Leung, "A survey on security threats and defensive techniques of machine learning," *IEEE Communications Surveys & Tutorials*, vol. 21, no. 1, pp. 765–794, First Quart. 2019.
- [10] J. Zhang, X. Chen, and Y. Xiang, "Federated learning: Privacy and security issues," *IEEE Internet of Things Journal*, vol. 8, no. 5, pp. 3566–3578, Mar. 2021.
- [11] A. Ghorbani, A. Abid, and J. Zou, "Interpretation of neural networks is fragile," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 45, no. 2, pp. 1969–1984, Feb. 2023.
- [12] S. Thrun and L. Pratt, "Learning to learn: Introduction and overview," *IEEE Transactions on Neural Networks*, vol. 10, no. 5, pp. 993–1000, Sept. 1999.
- [13] K. Chen, S. Patel, and R. Buyya, "Secure and scalable deployment of AI models in cloud environments," *IEEE Cloud Computing*, vol. 11, no. 1, pp. 34–45, Jan.–Feb. 2024.
- [14] V. Kamra, A. Singh, P. Sharma and R. Yadav, "A Novel Online Quizzing System for Blind People by implementing Modern Voice Recognition Techniques," 2023 2nd Edition of IEEE Delhi Section Flagship Conference (DELCON), Rajpura, India, 2023, pp. 1–4.
- [15] A. Müller and S. Guido, "Agent-based AI systems: Architecture and trust considerations," *IEEE Computer*, vol. 57, no. 3, pp. 48–57, Mar. 2024.
- [16] J. Kaur and M. Singh, "Towards standardized interfaces for trustworthy AI systems," *IEEE Transactions on Technology and Society*, vol. 5, no. 1, pp. 1–11, Mar. 2024.
- [17] T. Schick, H. Xu, and H. Ruder, "Toolformer: Language models can teach themselves to use tools," *arXiv preprint arXiv:2302.04761*, 2023.
- [18] T. Cai, X. Wang, and X. Chen, "Large language models as tool makers," *arXiv preprint arXiv:2303.12454*, 2023.
- [19] G. Feretzakis, K. Papaspyridis, and A. Gkoulalas-Divanis, "Privacy-preserving techniques in generative AI and large language models: A narrative review," *Information*, vol. 15, no. 11, p. 697, 2024.
- [20] K. Kalodanis, S. Papadopoulos, G. Feretzakis, K. Papaspyridis, and A. Gkoulalas-Divanis, "SecureLLM: A unified framework for privacy-focused large language models," *Applied Sciences*, vol. 15, no. 8, p. 4180, 2025.
- [21] V. Kamra, P. Kumar, and M. Mohammadian, "Diagnosis Support System for General Diseases by Implementing a Novel Machine Learning Based Classifier," *International Journal of Computing and Digital Systems*, vol. 10, no. 1, pp. 737–746, Nov. 2021.
- [22] S. Gaire, S. Gyawali, and S. Mishra, "Systematization of knowledge: Security and safety in the model context protocol ecosystem," *arXiv preprint arXiv:2512.08290*, 2025.
- [23] B. Radosevich and J. Halloran, "MCP safety audit: LLMs with the model context protocol allow major security exploits," *arXiv preprint arXiv:2504.03767*, 2025.